

```
In [1]: import medmnist
from medmnist import INFO
from torchvision import transforms
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.losses import BinaryCrossentropy
```

```
In [2]: # Load BreastMNIST dataset
data_flag = 'breastmnist'
download = True
info = INFO[data_flag]
DataClass = getattr(medmnist, info['python_class'])

# Define the transformations (resize to 128x128)
transform = transforms.Compose([
    transforms.Resize((128, 128)),
    transforms.ToTensor()
])

# Load training and test datasets
train_dataset = DataClass(split='train', transform=transform, download=download,
test_dataset = DataClass(split='test', transform=transform, download=download, s
```

```
In [3]: # Extract image data and labels (these will be NumPy arrays)
X_train = train_dataset.imgs
y_train = train_dataset.labels
X_test = test_dataset.imgs
y_test = test_dataset.labels

# Normalize and reshape the images
X_train = X_train.astype("float32") / 255.0
X_test = X_test.astype("float32") / 255.0

# Add the channel dimension: (N, H, W, C) -> (N, 128, 128, 1)
X_train = np.expand_dims(X_train, axis=-1) # Shape: (N, 128, 128, 1)
X_test = np.expand_dims(X_test, axis=-1)   # Shape: (N, 128, 128, 1)


# Checking the shapes
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
```

X_train shape: (546, 128, 128, 1)


X_test shape: (156, 128, 128, 1)

```
In [4]: # Display first 5 training images
for i in range(5):
    plt.imshow(X_train[i].squeeze(), cmap='gray') # .squeeze() to remove the ch
    plt.title(f"Label: {y_train[i]}")
    plt.axis('off')
    plt.show()
```


Label: [1]




Label: [1]




Label: [1]



Label: [1]



Label: [0]



```
In [5]: # Build the Keras CNN model
model = Sequential([
    Input(shape=(128, 128, 1)), # Define the input shape (128x128 images with 1
    Conv2D(32, (3,3), activation='relu'),
    MaxPooling2D(pool_size=(2,2)),

    Conv2D(64, (3,3), activation='relu'),
    MaxPooling2D(pool_size=(2,2)),

    Conv2D(128, (3,3), activation='relu'),
    MaxPooling2D(pool_size=(2,2)),

    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.5),
    Dense(1, activation='sigmoid') # Output layer for binary classification
])

# Compile the model
model.compile(optimizer=Adam(),
              loss=BinaryCrossentropy(),
              metrics=['accuracy'])

# Model summary
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 126, 126, 32)	320
max_pooling2d (MaxPooling2D)	(None, 63, 63, 32)	0
conv2d_1 (Conv2D)	(None, 61, 61, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 30, 30, 64)	0
conv2d_2 (Conv2D)	(None, 28, 28, 128)	73,856
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 128)	0
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 128)	3,211,392
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 1)	129



Total params: 3,304,193 (12.60 MB)

Trainable params: 3,304,193 (12.60 MB)

Non-trainable params: 0 (0.00 B)

```
In [6]: # Train the model
model.fit(X_train, y_train, epochs=10, batch_size=64, validation_split=0.1)
```

```

Epoch 1/10
8/8 ————— 3s 257ms/step - accuracy: 0.6601 - loss: 0.6689 - val_ac
curacy: 0.7636 - val_loss: 0.5084
Epoch 2/10
8/8 ————— 2s 232ms/step - accuracy: 0.7600 - loss: 0.5437 - val_ac
curacy: 0.7636 - val_loss: 0.5705
Epoch 3/10
8/8 ————— 2s 225ms/step - accuracy: 0.7168 - loss: 0.5616 - val_ac
curacy: 0.7636 - val_loss: 0.4953
Epoch 4/10
8/8 ————— 2s 217ms/step - accuracy: 0.7226 - loss: 0.5625 - val_ac
curacy: 0.7636 - val_loss: 0.5179
Epoch 5/10
8/8 ————— 2s 226ms/step - accuracy: 0.7427 - loss: 0.5038 - val_ac
curacy: 0.7818 - val_loss: 0.5051
Epoch 6/10
8/8 ————— 2s 238ms/step - accuracy: 0.7333 - loss: 0.5153 - val_ac
curacy: 0.7636 - val_loss: 0.5012
Epoch 7/10
8/8 ————— 2s 218ms/step - accuracy: 0.7323 - loss: 0.5180 - val_ac
curacy: 0.8000 - val_loss: 0.5117
Epoch 8/10
8/8 ————— 2s 226ms/step - accuracy: 0.7755 - loss: 0.4766 - val_ac
curacy: 0.7818 - val_loss: 0.4967
Epoch 9/10
8/8 ————— 2s 242ms/step - accuracy: 0.7394 - loss: 0.5006 - val_ac
curacy: 0.8000 - val_loss: 0.4850
Epoch 10/10
8/8 ————— 2s 236ms/step - accuracy: 0.7813 - loss: 0.4383 - val_ac
curacy: 0.7636 - val_loss: 0.4763

```

Out[6]: <keras.src.callbacks.history.History at 0x22e7cc8bd70>

```

In [7]: # Evaluate the model on test data
loss, accuracy = model.evaluate(X_test, y_test)
print(f"Test Accuracy: {accuracy:.4f}")

```

```

5/5 ————— 0s 36ms/step - accuracy: 0.7642 - loss: 0.4693
Test Accuracy: 0.7756

```

```

In [8]: # Make predictions on the test data
predictions = (model.predict(X_test[:5]) > 0.5).astype("int32")

# Display the first 5 predictions
for i in range(5):
    plt.imshow(X_test[i].reshape(128,128), cmap='gray')
    plt.title(f"Pred: {predictions[i][0]} | Actual: {y_test[i][0]}")
    plt.axis('off')
    plt.show()


```

```


1/1 ————— 0s 92ms/step

```

Pred: 1 | Actual: 0




Pred: 1 | Actual: 1




Pred: 1 | Actual: 1



Pred: 1 | Actual: 1



Pred: 1 | Actual: 1



In []: